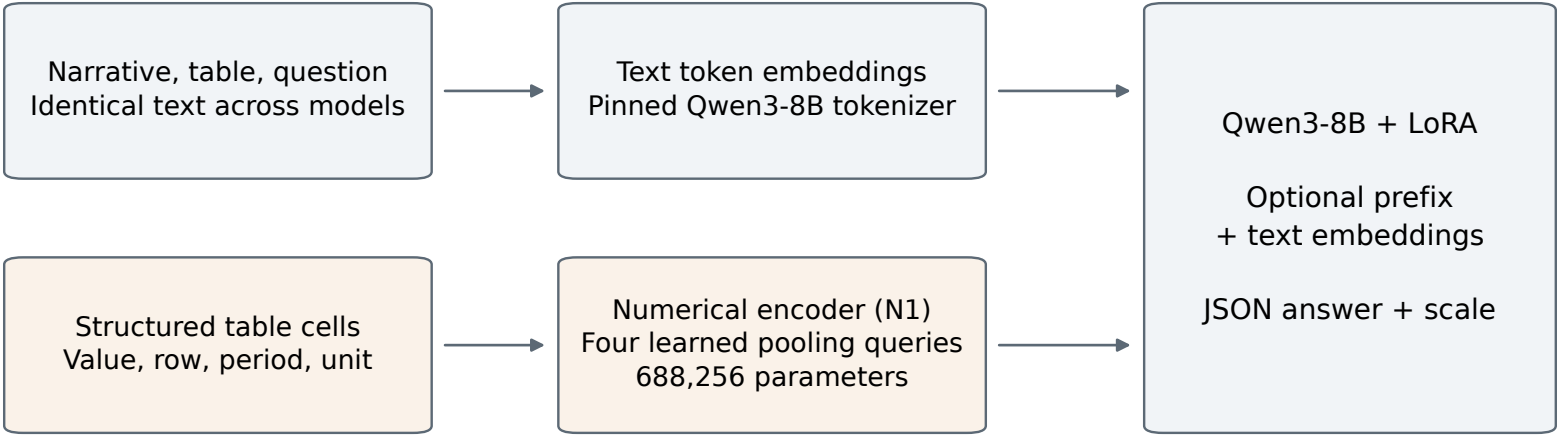
Narrative, table, question
Identical text across models
Text token embeddings
Pinned Qwen3-8B tokenizer
Qwen3-8B + LoRA
Optional prefix
+ text embeddings
JSON answer + scale
Structured table cells
Value, row, period, unit
Numerical encoder (N1)
Four learned pooling queries
688,256 parameters
B1: no prefix | N1: four cell-conditioned tokens | P1: four input-independent tokens