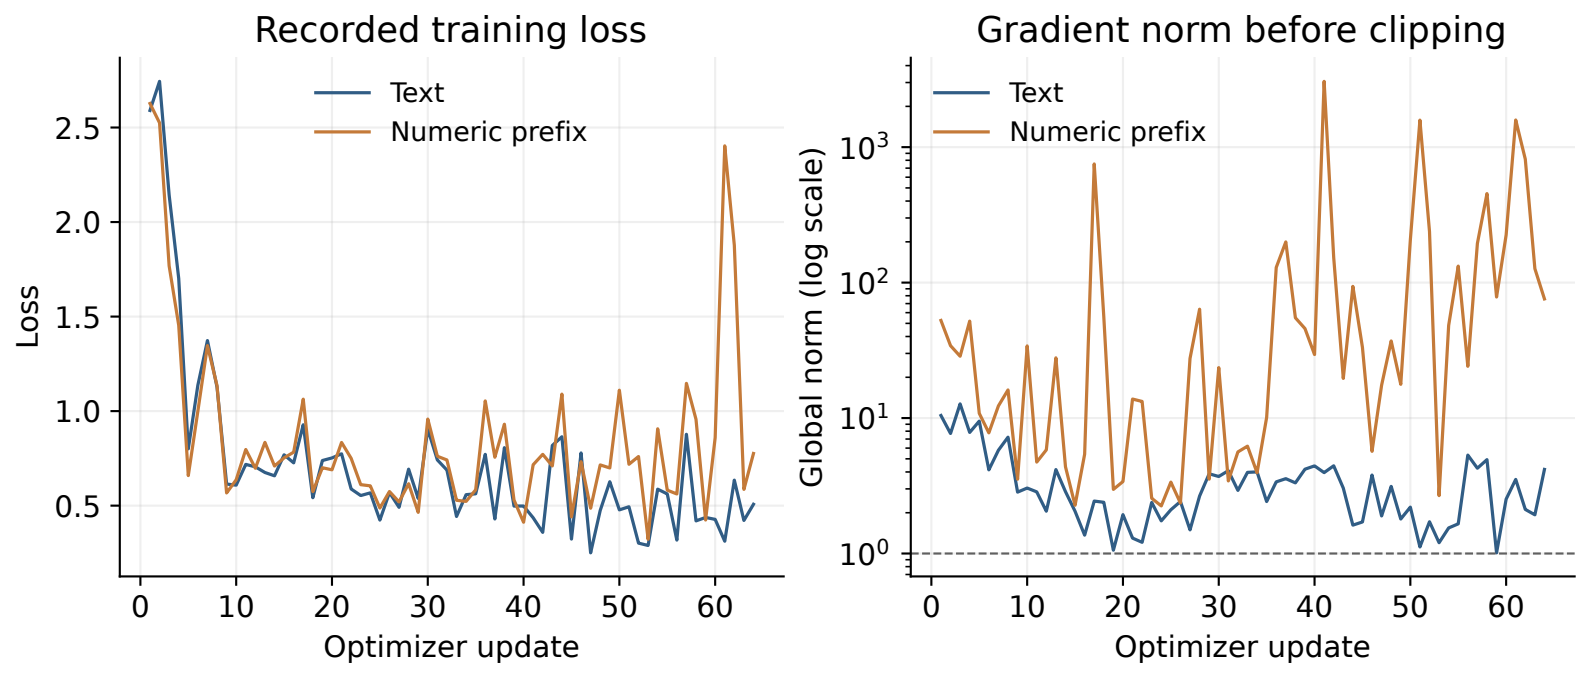
Recorded training loss
Text
Numeric prefix
Loss
2.5
2.0
1.5
1.0
0.5
0
10
20
30
40
50
60
Optimizer update
Gradient norm before clipping
Text
Numeric prefix
Global norm (log scale)
10^3
10^2
10^1
10^0
0
10
20
30
40
50
60
Optimizer update